

Architecture-Specific Evaluation of LLM-Generated Code Optimizations on x86 CPUs

Hirotaka Monya¹ Taro Watanabe¹ Sakriani Sakti¹ Hidetaka Kamigaito¹ Yusuke Sakai¹

¹ Nara Institute of Science and Technology NLP Laboratory

Motivation / Background

- C optimization is arch-dependent and labor-intensive
- LLMs can follow textual design rules; architecture info can be prompted
- Need a safe loop: generate quickly but verify performance & correctness

Objective

- Quantify when LLM-made C beats gcc -O3 per architecture
- Map effective transforms to AVX/AVX2/AVX-512
- Provide a reproducible protocol; release prompts/variants

Related Work

- Auto-tuning/DSL/Superopt: Ansor, Exo/Exocompilation, Mirage
- LLM-centric: Self-Refine, LLM Compiler; program-repair survey

Experimental Setup

- PolyBench-C (30 kernels)
- LLMs: GPT / Claude (n-best)
- CPUs: snb-avx, hsw-avx2, skx-avx512, icl-avx512
- Baselines: gcc -O0/-O3; pinning; median-of-7

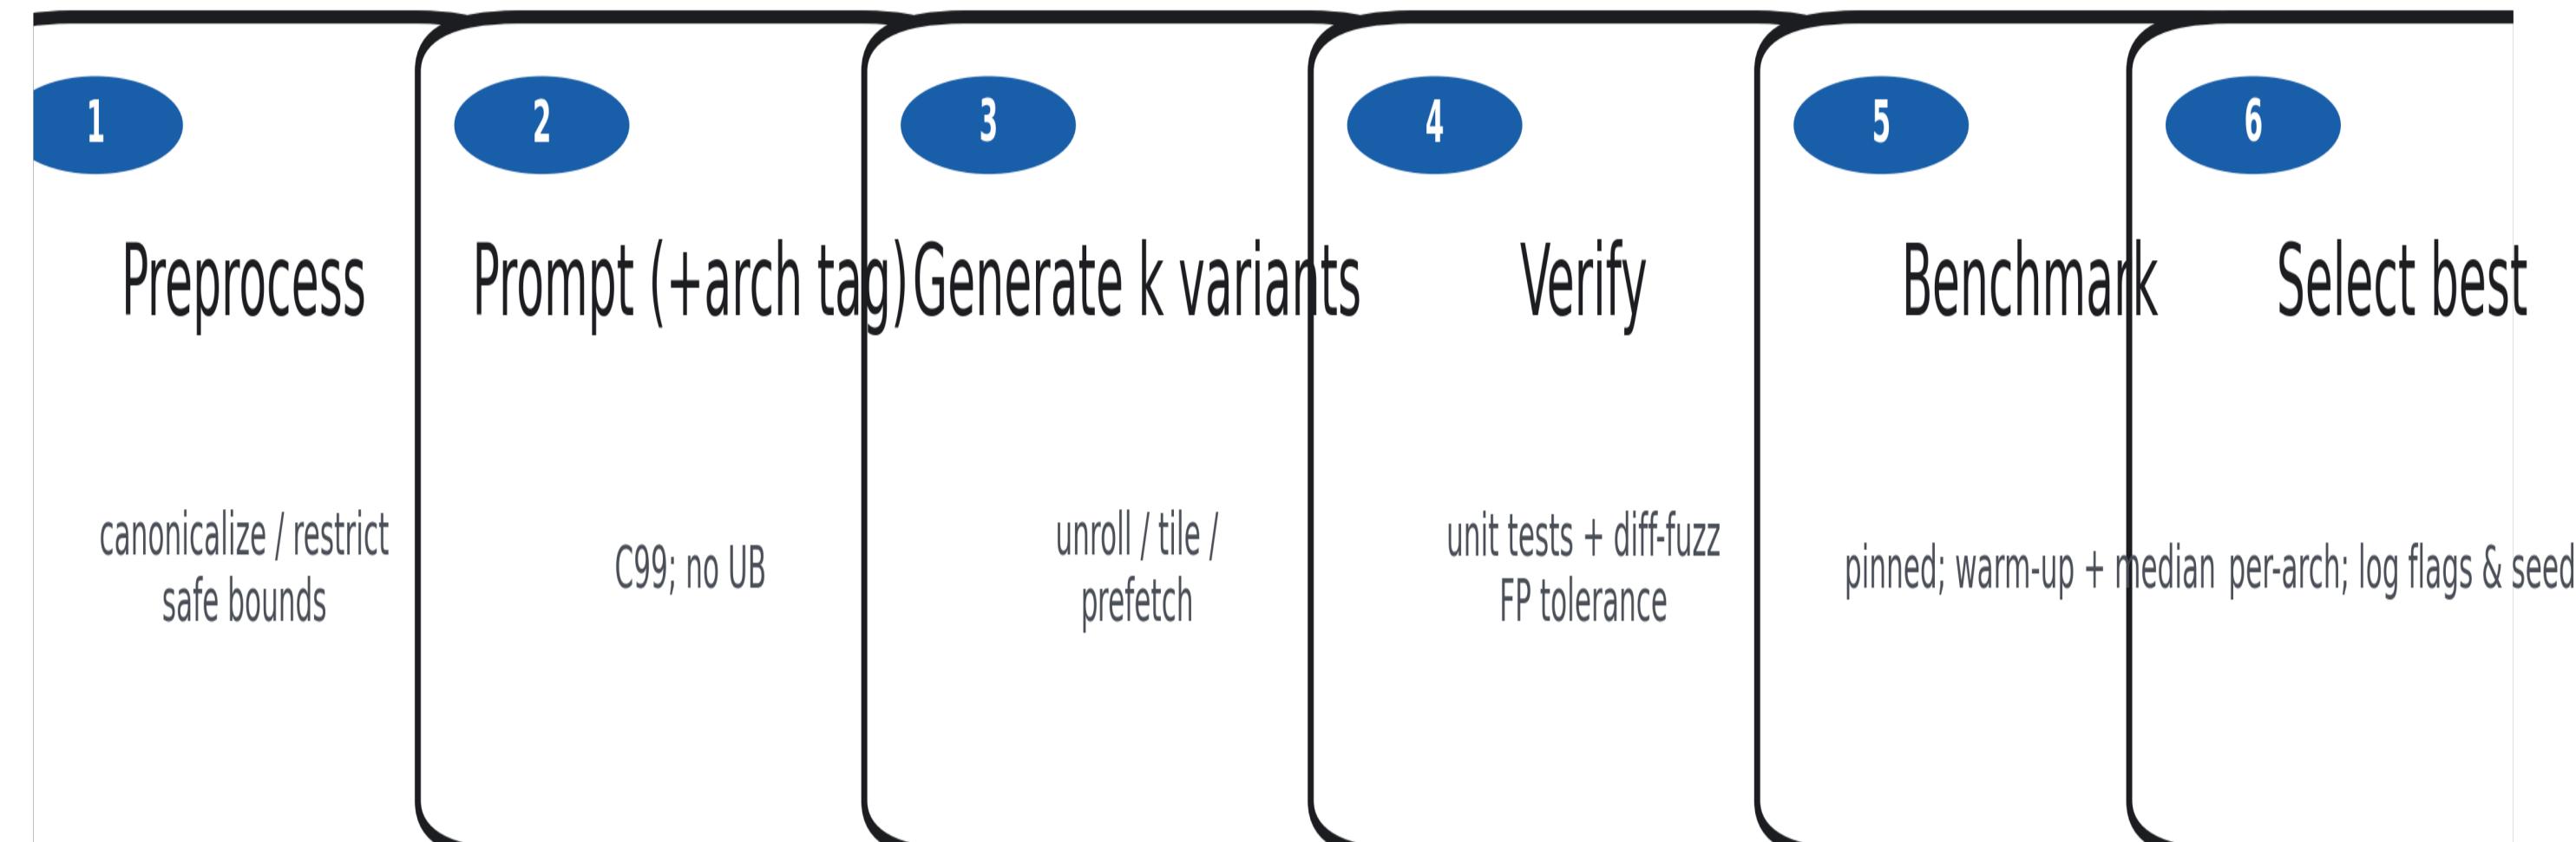
Results II: Kernel-wise Best

| Kernel | Category | Best arch | Best LLM | Speedup vs -O3 | Key transform |
|-----------|-----------|------------|----------|----------------|---------------------------------|
| gemm | BLAS-like | skx-avx512 | GPT | 1.48× | tile 64×64; prefetch 64; FMA |
| 2mm | BLAS-like | skx-avx512 | GPT | 2.05× | tiled + unrolled; aligned loads |
| 3mm | BLAS-like | icl-avx512 | GPT | 1.62× | tile 32×32; prefetch 128 |
| syrk | BLAS-like | hsw-avx2 | Claude | 1.36× | unroll ×4; restrict |
| atax | Regular | hsw-avx2 | Claude | 1.22× | vectorize + unroll; bounds fix |
| bicg | Regular | snb-avx | GPT | 1.15× | loop reorder; align |
| fdtd-2d | Stencil | icl-avx512 | GPT | 1.28× | tile 32×32; mask tails |
| jacobi-2d | Stencil | hsw-avx2 | GPT | 1.18× | unroll ×4; prefetch 64 |

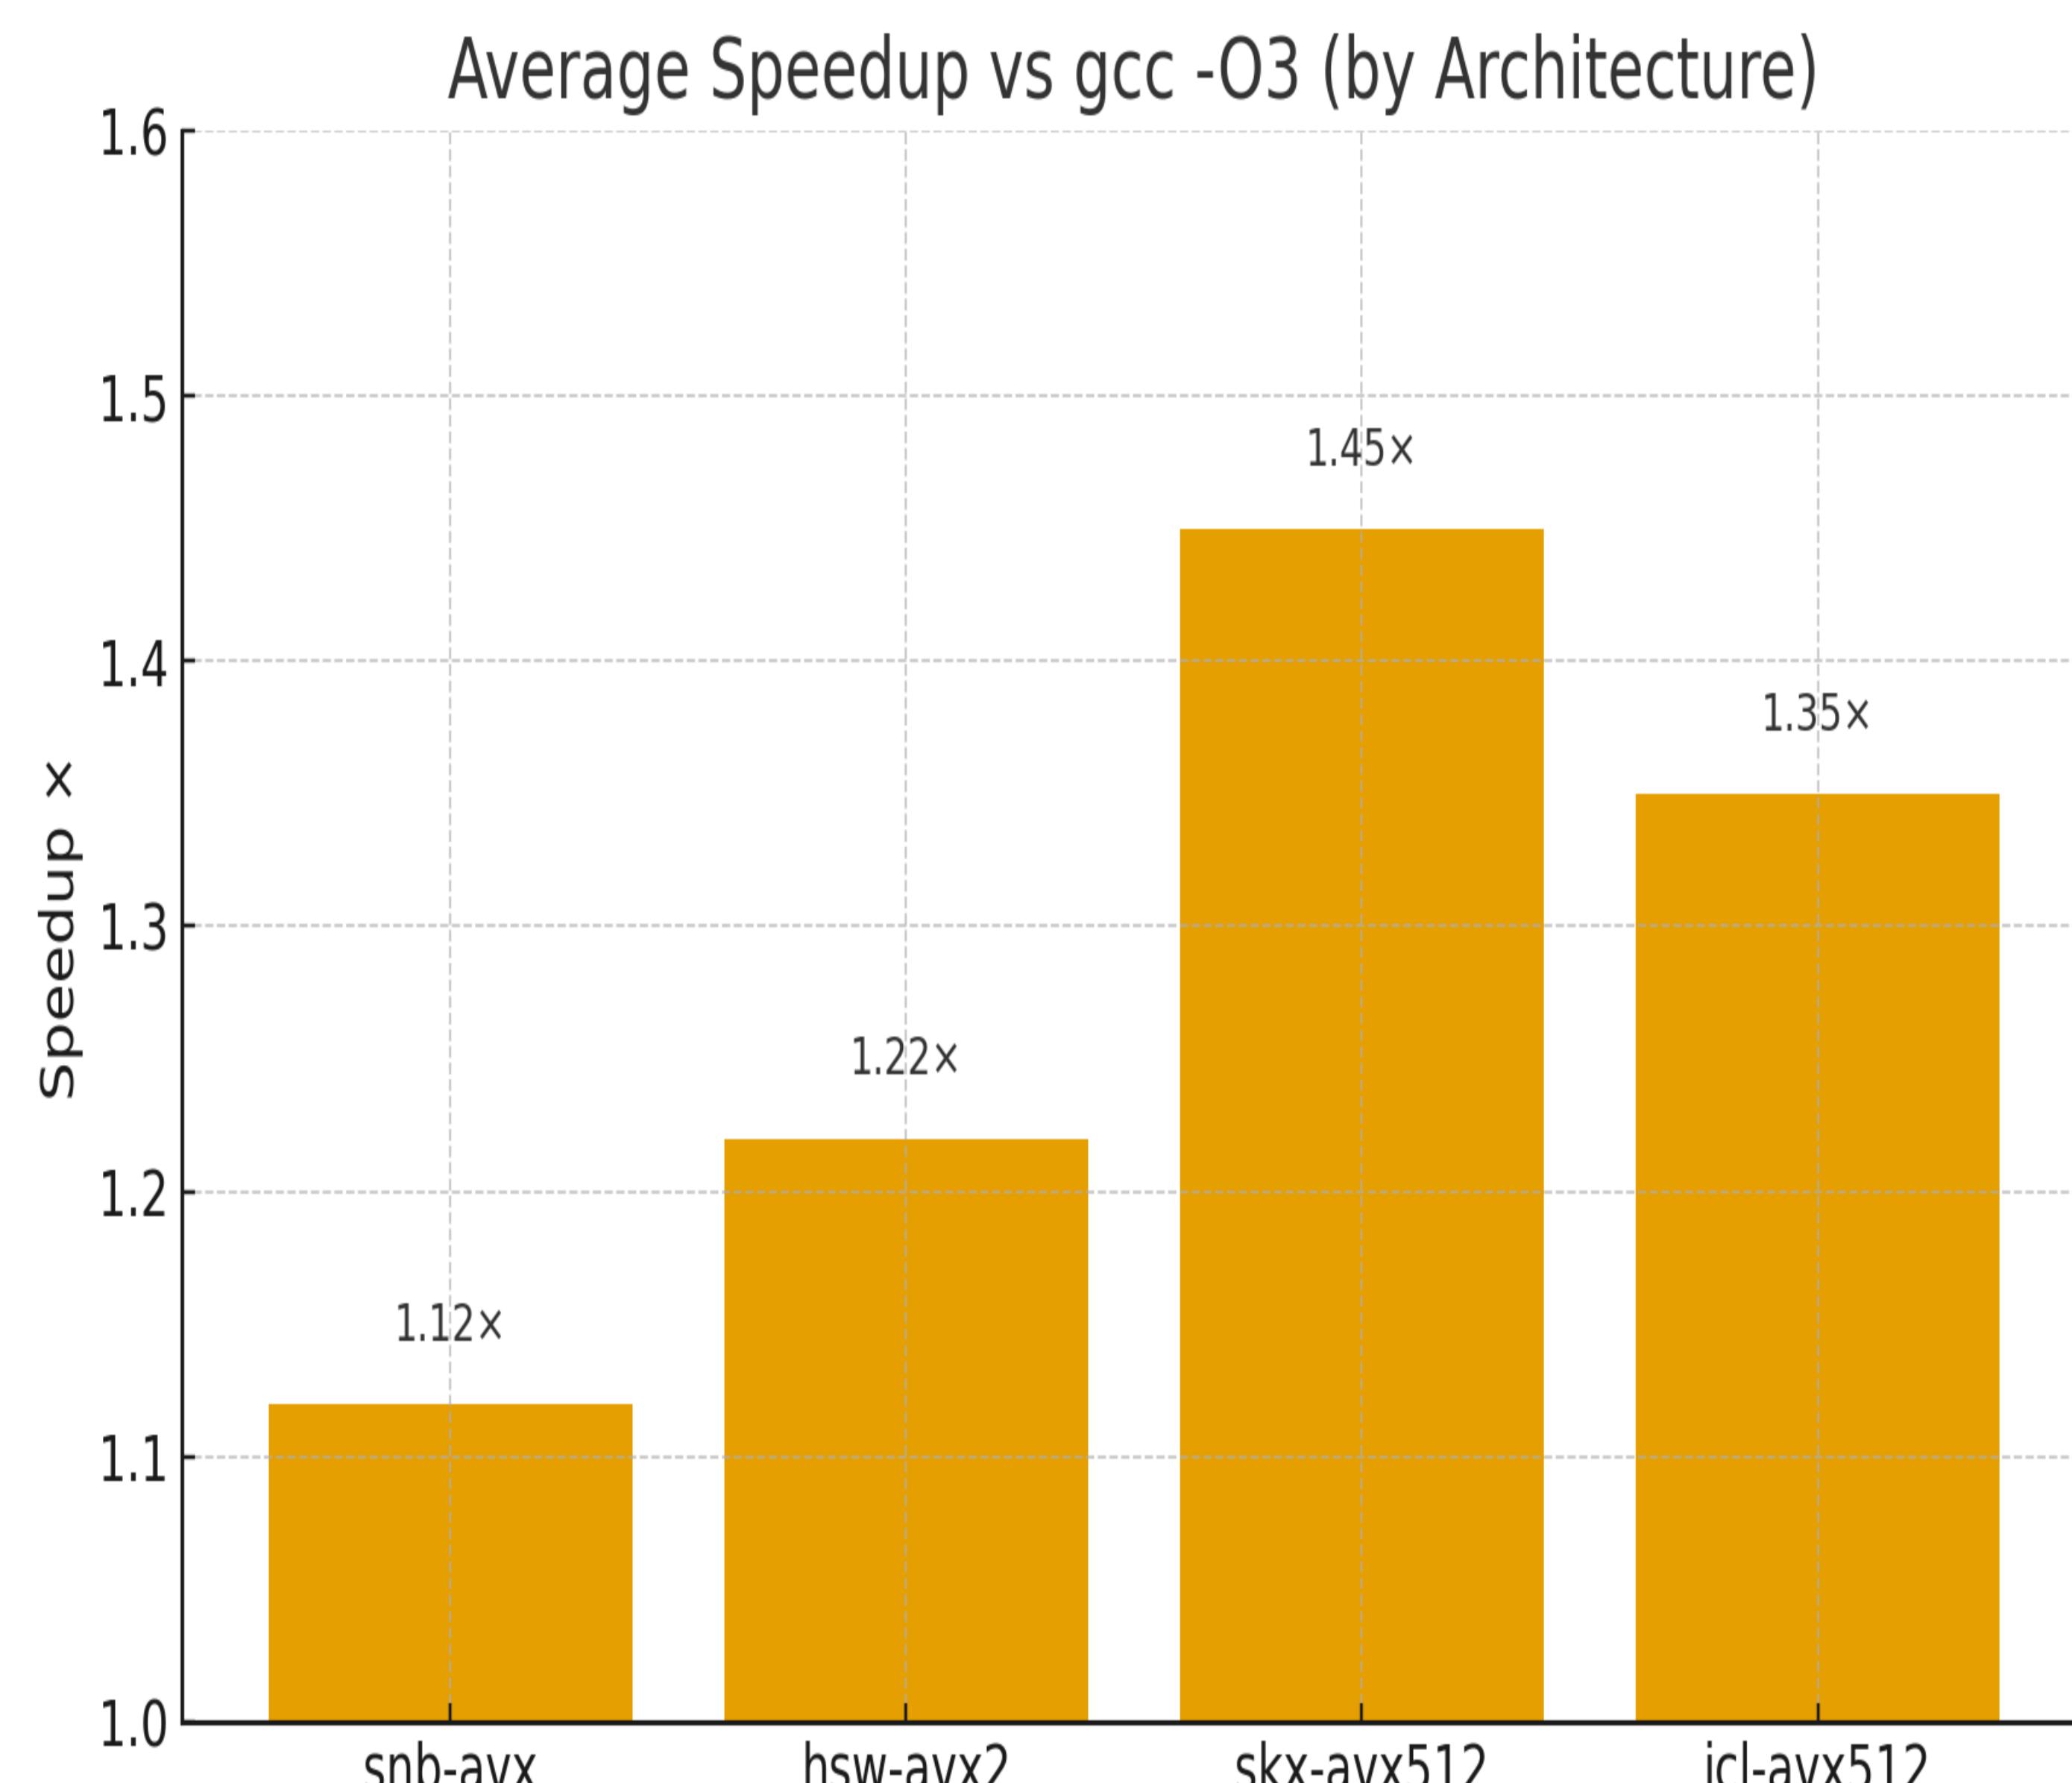
Optimizations by Architecture

- snb-avx: small tiles, unroll×2, aligned contiguous loads
- hsw-avx2: unroll×4, 32–64 tiles, FMA-friendly schedule
- skx-avx512: 32–64 tiles, masked tails, mid-long prefetch
- icl-avx512: 32 tiles, short prefetch, keep FMA density high

Method (Closed Loop)



Results I: Average Speedup by Architecture



Conclusion & Next Steps

Architecture-tagged prompts + verify-then-benchmark improve C performance safely. Next: profile-guided schedules, light auto-tuning, stronger equivalence (SMT/Lean), and GPU/Triton.